

YAHOO!

Offense At Scale

PRESENTED BY Chris Rohlf | May 30th, 2015

whoami

- I lead the pentesting team at [Yahoo!](#)
- Founded Leaf Security Research, a boutique security consultancy
- Previously {Principal Consultant Matasano Security, DoD}
- Black Hat review board member and speaker (2009, 2011, 2013)
- Security researcher and tool/exploit developer

Introduction

This keynote

Introduction

What is a system in 2015?

Introduction


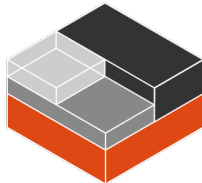
What is scale?

Disposable Environments

"Program against your datacenter like it's a single pool of resources"



Container, Unikernel, Picoprocess,
Hypervisors



YAHOO!

Offense and Scale

- We can break down offense into three areas
 - Vulnerability Research, Pentesting, Red Teaming
- How does the offense view architectures designed to scale?
 - The bottom of the stack moves slow
 - The top of the stack changes fast
 - Reliability is a transitive property
 - We can't compete without automation

Offense and Scale

- The massive growth of infrastructures has resulted in the unexpected
- We rely too much on the 1 in a million defense
 - “In our business, one in a million is next Tuesday”
 - <http://blogs.msdn.com/b/larryosterman/archive/2004/03/30/104165.aspx>
- “Mean time between failure doesn’t scale” - @iamreallyfrank

Offense and Scale: Scenario 1

- You want to gain access to as many email accounts as possible
- Steal/buy a password list, take the most used password and try it one time against 1 billion email accounts
 - Low risk of massive account lockouts
 - Almost guaranteed to have a decent success rate

Offense and Scale: Scenario 2

- A forking HTTPD with an exploitable heap overflow
 - x64 Linux ASLR Heap has 28 bits of entropy
 - 2^{28} possible base address values (268,435,456)
 - $268,435,456 / 200,000$ httpd daemons = 1342 requests
 - $268,435,456 / 300,000$ httpd daemons = 894 requests
 - $268,435,456 / 400,000$ httpd daemons = 671 requests
 - $268,435,456 / 500,000$ httpd daemons = 536 requests
 - Hint: 1342 requests per second was impressive in mid-2000s
 - The birthday problem applies to ASLR
 - Can you analyze a terabyte of logs and respond before the entire file system is exfiltrated?

Offense and Scale

- *“Heartbleed Team, National Security Agency. Their efforts resulted in a means to stop attempts to exploit a security vulnerability across the department’s **global network of more than eight million computing devices.**”*
 - <http://www.defense.gov/releases/release.aspx?releaseid=17062>
- *“**The exercise took place on a specially-constructed closed network designed to simulate the DoD and allied information networks and adversary networks.** The event also featured an expert opposing force, which takes on the role of the adversary, using a range of tactics and weapons to provide a realistic training environment.”*
 - <http://www.defense.gov/news/newsarticle.aspx?id=123621>
 - How can you simulate the real world scale (e.g. internet) these people operate in?
- *“The United States Department of Defense (DoD) Cyber Range is a **realistic simulation** and modeling network environment”*
 - <http://www.tintri.com/customers/departement-defense-cyber-range>

Vulnerability Research at Scale

- Targeting
 - Ubiquitous software is a good start
- Data driven analysis
 - Where do you spend your limited resources and time?
 - Continuously consume and analyze vulnerability data
- Fuzzing
 - Containers can increase scale for certain applications
 - yFuzz - Automated fuzzing at scale with Docker
- Automation
 - Too much of our work is still a manual process

Offensive tools are not keeping up

- Most public tools are not built to scale
 - No concept of one-to-many
 - Often built for and tested against 1 to 5 systems
 - Never able to handle more than STDOUT logging
 - Rely too heavily on a single language (Python)
 - Worms written in 2001 are more capable of lateral movement than most pentesters
- Majority of focus is on client side

Pentesting At Scale

- Vulnerability research has paid off, you have 0day to deploy to thousands of systems
 - Revisiting your tools is the next logical step
- Scenario: You have XXE 0day on a popular web framework and the network you are targeting has it widely deployed
 - Your favorite tool *abandoned_web_scanner.py* will take 4 days to complete
 - It has to perform a DNS lookup for each host and has no async resolver
 - The tool has no concept of threads, and even if it did Python has a GIL
 - My solution at Yahoo was to wrap libcurl with a thread pool library in C++
 - 10x as fast as the earlier Ruby implementation
 - Took no more than a week to develop
 - Can be extended to create many other web scanners



**FIFA Officials
Charged in
Corruption Probe**



**New Push Ties Cost
of Drugs to How Well
They Work**



**Companies Send
More Cash Back to
Shareholders**



**BOE Official
Received Emails
Relating to Libor
Manipulation**



**Brexit
Exposures**



25



10


LEADER

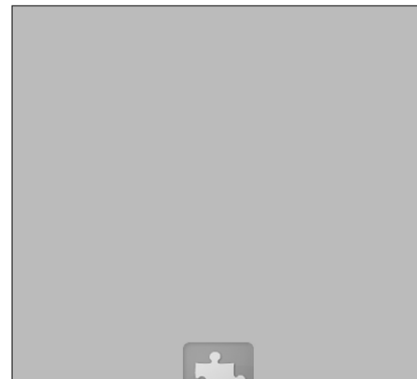
Of Microbes and Mock Attacks: Years Ago, The Military Sprayed Germs on U.S. Cities

By **JIM CARLTON** Staff Reporter of The Wall Street Journal

Updated Oct. 22, 2001 12:01 a.m. ET

SAN FRANCISCO -- Fifty-one years ago, Edward J. Nevin checked into a San Francisco hospital, complaining of chills, fever and general malaise. Three weeks later, the 75-year-old retired pipe fitter was dead, the victim of what doctors said was an infection of the bacterium *Serratia marcescens*.

Decades later, Mr. Nevin's family learned what they believe was the cause of the infection, linked at the time to the hospitalizations of 10 other patients. In Senate subcommittee hearings in 1977, the U.S. Army revealed that weeks...



Discovering attack surface

Physical



Code

The screenshot displays the Ghidra decompiler interface with the following components:

- Function Header:**

```

Attributes: bp-based frame
asdlContext: clientHeaderType(unsigned int, assigned int *, IDHeaderDescriptor **)
__initContextVcContextHeaderType(p91HeaderDescriptor *, proc_ptr)
var_30: guard_ptr - 30h
var_30: guard_ptr - 30h

```
- Assembly Code:**

```

push rbp
mov rbp, rbp
push r15
push r14
push r13
push r12
push r11
sub rsp, 10h
mov r14, r14
mov r12, r12
mov r15, r15
mov eax, [r15+50h]
lock
mov eax, eax
je short loc_202225

```
- Control Flow Graph (CFG):**
 - Block 1 (loc_202225):**

```

loc_202225:
[rsp, r14, 5]
inc loc_202200

```
 - Block 2 (loc_202200):**

```

loc_202200:
mov rax, cs:off_09130
mov r14, r15
mov r12, r12
mov r13, r14
and r13, 10h
pop r11
pop r10
pop r13
pop r15
mov r10, r10
call qword ptr [rax+210h]
je short loc_20227F

```
 - Block 3 (loc_20227F):**

```

loc_20227F:
mov r10, [r15+100h]
mov [rbp+30h], r10
mov r14, [rbp+30h]
call [rbp+30h]
mov r14, [rbp]
mov rax, [rbp+graphicsAccelerator2Witch_busyq_10graphicsAccelerator2Lock_busyq]
lea r10, off_50322
xor edx, edx
mov r14, r10
call qword ptr [rax+510h]
mov r10, [r15+100h]
mov r14, [r15+100h]
test r10, r10
je short loc_20227F

```

Arrows indicate the control flow between these blocks, showing a loop structure.

Network/Data



Red Teaming At Scale

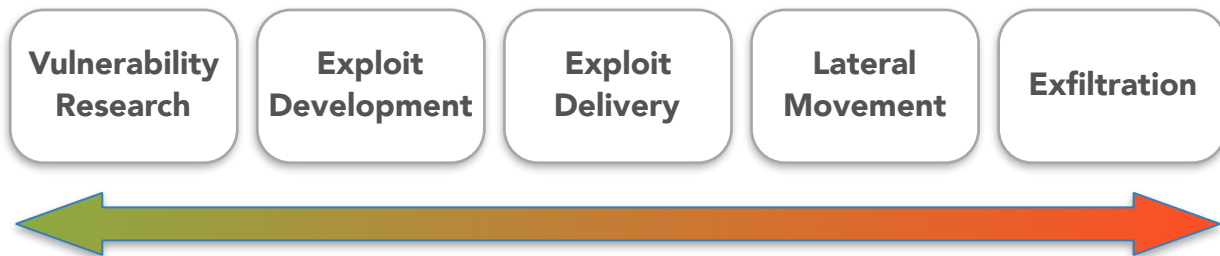
- An effective red team (and defenders) studies the real world tools, techniques and procedures from past attack campaigns
 - What about your attackers methods are different because of the size and scale of your operation?
 - Did their method of exfiltration take advantage of speed or could you have beaten them had you known what to look for?
 - Did the growth of the haystack make it harder to find the needle?
 - Does the maturity of your operation lend itself to faster needle searching?



Attack Chains


- Scale can afford you the option of multiple paths to a target
 - This is an opportunity to avoid overlapping IOCs and toolchains
- Crash and/or evasion of security products
 - If defenders are using speed to their advantage in log collection and analysis then identify and exploit the slowest part of that process
 - NIDS are easily bypassed and taken down via memory consumption, NULL pointer dereferences and volume of noise
 - Many of these technologies fail open by design
- Continuous Integration / Continuous Delivery
 - Built in lateral movement

Automation

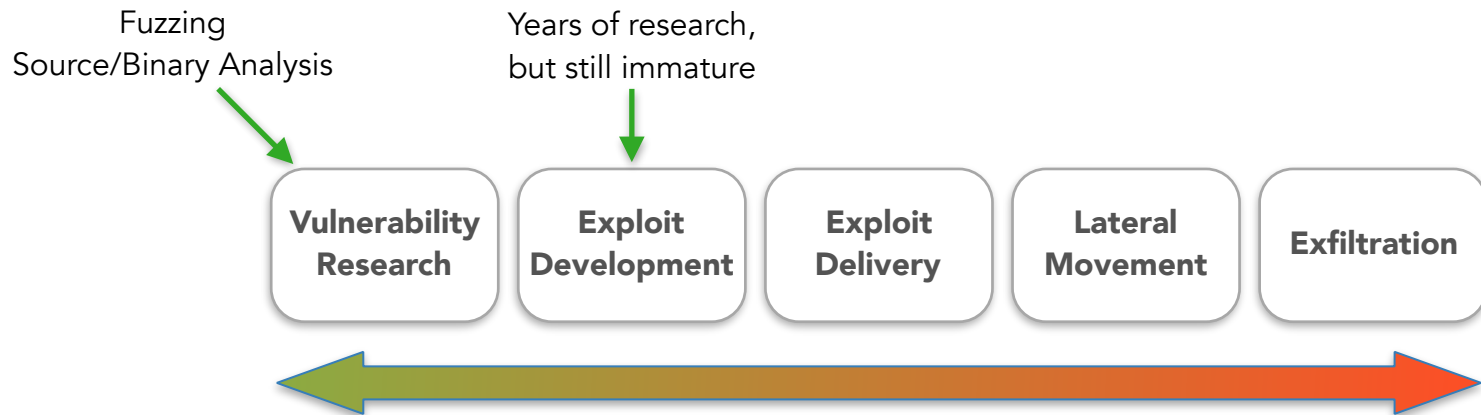


Automation

Fuzzing
Source/Binary Analysis



Automation




Automation

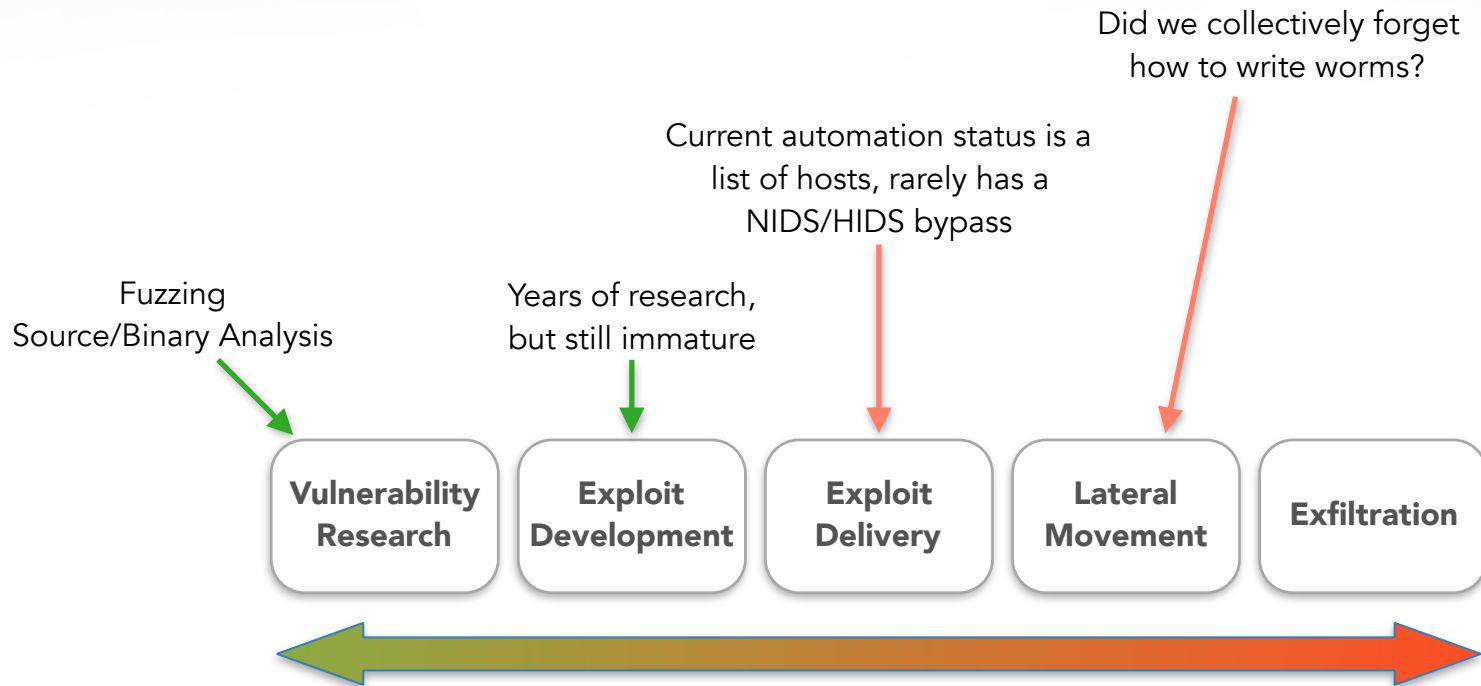
Current automation status is a
list of hosts, rarely has a
NIDS/HIDS bypass

Fuzzing
Source/Binary Analysis

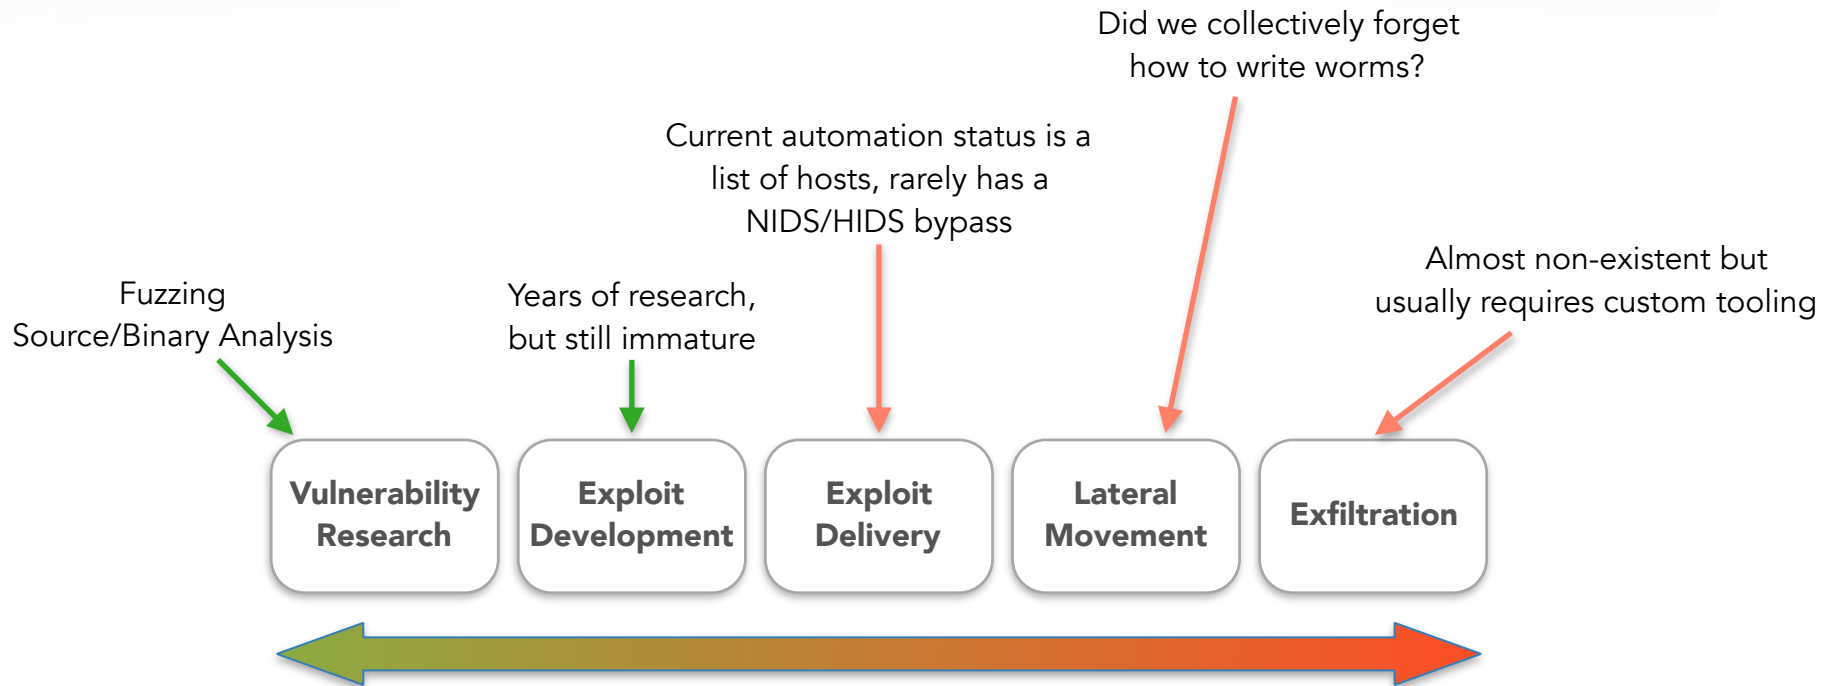
Years of research,
but still immature



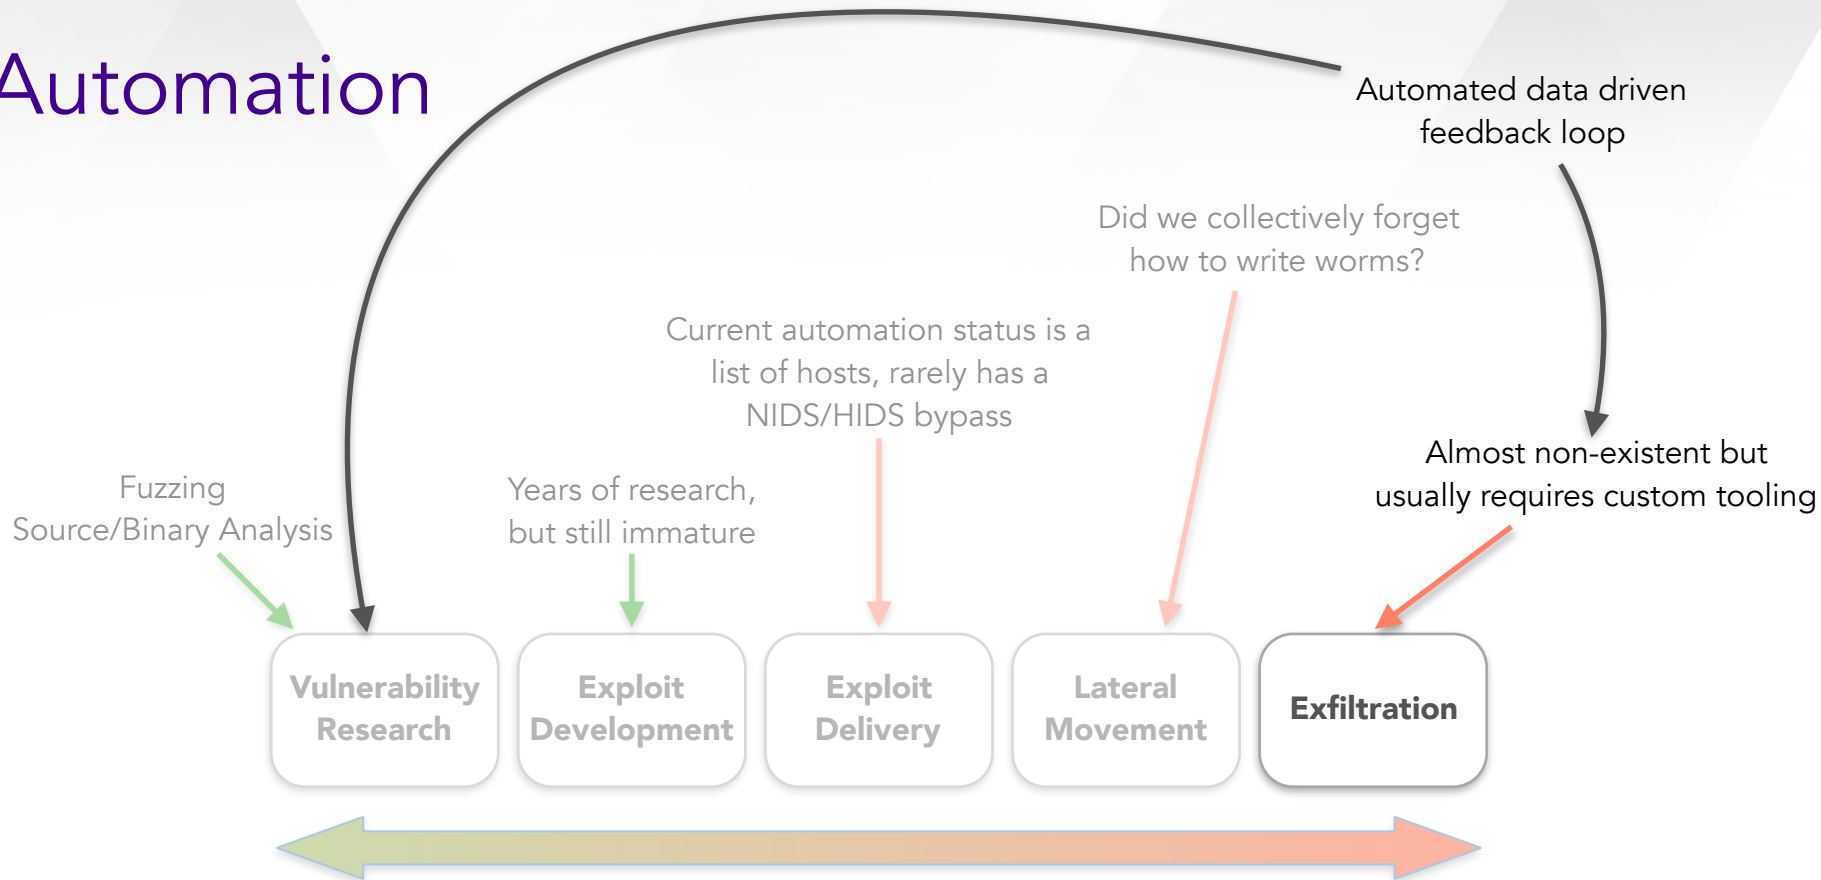
Automation



Automation



Automation



Post Compromise

- The 'Low and Slow' approach is out
- Exfiltration will have to meet or exceed speed of log collection/analysis
- There is no HIDS with a knowledge of containers
 - Containers have a shared kernel
 - Until Docker meets SECCOMP-BPF, the bar for compromising another container is lower than a hypervisor
 - HIDS will need some time to catch up

The Future Of Disposable Environments

- Hypervisors, Containers, Unikernels, Picoprocesses
 - Scenario: It's 2020 and every web server request is handled by a picoprocess on demand that is terminated upon TCP FIN
 - Exfiltration must be automated, over TLS and faster than log analysis can detect it
 - Can your current toolset do this?
 - Can your current toolset catch this?
- Vulnerabilities that allow an attacker to escape these environments are valuable
 - Row Hammer
 - Hypervisor memory corruption (<http://xenbits.xen.org/xsa>)

Advice for defenders

- Hardened targets with dense bugs are still more secure than targets with no hardening and sparse bugs (e.g. Chrome vs Bash)
- A moving target is always harder to exploit
 - Reduces lifetime and increases cost of an exploit
 - CI/CD can be leveraged to achieve this
- Security staff levels
 - %1 of total headcount
 - Security staff that focuses on enterprise systems should be measured against system count
 - Security staff focused on product should be measured by application complexity

Advice for defenders

- Attacker asymmetry, like everything else, grows linearly with scale
- Scale has a way of magnifying even the smallest security problems
 - “1 in a million is next Tuesday”
- Work with what you’ve got and put pressure on your vendors to deliver secure software
- The future of defending scale is $O(n)$ where n is a list of systems
 - Is your *detect_respond()* faster than *identify_pwn()* ?

The End

Scale? ^H^H^H^H Questions?